

BSC INF Course - Graph-Partitioning

October 18, 2022

Olaf Schenk
Institute of Computing
Faculty of Informatics
USI Lugano

Content


- Motivation for graph partitioning
- Overview of heuristics
- Partitioning with nodal coordinates
 - Ex: In finite element models, node at point in (x, y, z) space
 - Recursive Coordinate Bisection**
 - Inertial Partitioning**
- Partitioning without nodal coordinates
 - Ex: In model of WWW, nodes are web pages
 - Spectral Methods**

Partitioning and Load Balancing


- **Goal:** assign data to processors to
 - minimize parallel application runtime
 - maximize utilization of computing resources
- **Metrics:**
 - minimize processor idle time (balance workloads)
 - keep inter-processor communication costs low
- Impacts performance of a wide range of simulations

A 6x6 grid representing matrix A, with red squares indicating non-zero entries. To its right is a 6x1 column of white squares representing vector x, followed by an equals sign, and then a 6x1 column of pink squares representing vector b.


Linear solvers & preconditioners



Contact detection



Particle simulations




Adaptive mesh refinement

Graph Partitioning

- Work-horse of load-balancing community.
- **Highly successful model for PDE problems.**
- Model problem as a graph:
 - vertices = work associated with data (computation)
 - edges = relationships between data/computation (communication)
- Goal: Evenly distribute vertex weight while minimizing weight of cut edges.

Definition of Graph


- Given a graph $G = (V, E)$ with
 - Vertices $V = \{ v_i \mid i=1,\dots,n \}$
 - Edges $E = \{ e_{ij} \mid v_i \text{ and } v_j \text{ are connected} \}$



$$V = \{ 1, 2, 3, 4 \}$$

$$E = \{ (1,2), (1,3), (2,3), (2,4) \}$$

- A weighted graph $G = (V, E, W_v, W_e)$ has **node weights** and **edge weights**
 - $W_v = \{ w_v(v_i) \mid v_i \in V \}$ („weight of vertices“).
 - $W_e = \{ w_e(e_{ij}) \mid e_{ij} \in E \}$ („weight of edges“).




$$W_v = \{ 3, 3, 7, 15 \}, W_e = \{ 4, 5, 8, 1 \}$$

Examples for Graphs

- Symmetric sparse matrix and Graph G_A

$A =$

	1	2	3	4	5	6
1	1	1			7	
2	8	6		1		1
3			3	1		1
4		5	1	2	1	
5	2			1	5	1
6		1	5		1	1



- $G_A = (V, E, W_v, W_e); V = \{1, 2, 3, 4, 5, 6\},$



$$E = \{ (1,2), (1,5), (2,4), (2,6), (3,4), (3,6), (4,5), (5,6) \}$$

$$W_v = \{2, 3, 2, 3, 3, 3\} \text{ e.g. numbers of nonzeros in each row}$$


$$W_e = \{1, 1, 1, 1, 1, 1, 1, 1\}$$

Examples for Graphs

- Finite-Element Simulations



Finite-Element Mesh






$$G_{FE} = (V, E), V = \{1, \dots, 14\}$$

$$E = \{(1,2), \dots, (12,14)\}$$

$$W_e \equiv 1, W_v \equiv 1$$

Examples for Graph Partitioning

- Parallel Finite-Element Simulations




A good partitioning G_{FE} results in

- equal #elements/processor („load“ and „storage balancing“).
- Minimal #edges between P1 and P2 (minimal communication volume).

Definition of Graph Partitioning: Bisection

- Given a graph $G = (V, E, W_V, W_E)$
 - V = nodes (or vertices)
 - E = edges



- Choose a partition $V = V_1 \cup V_2$ such that:
The sum of the node in each V_j is “about the same”

$$V = V_1 \cup V_2, \quad V_1 \cap V_2 = \emptyset, \quad |V_1| = |V_2|$$


The sum of edge connecting pairs V_1 and V_2 is minimized

$$\min |\{e_{ij} \in E \mid v_i \in V_1 \text{ und } v_j \in V_2\}|$$

Heuristics — Overview of Bisection Algorithms

- Partitioning with nodal coordinates — e.g. each node has x,y,z coordinates → partition space

Algorithms: **Recursive Coordinate Bisection**
Inertial Partitioning




- Partitioning without Nodal Coordinates — e.g. indexing of web documents $A(j,k) = \# \text{ times keyword } j \text{ appears in URL } k$

Algorithms: **Spectral Methods**

Nodal Coordinates — Recursive Coordinate Bisection (RCB)


- Developed by Berger & Bokhari (1987)
 - Independently discovered by others.

- Idea:
 - Divide work into two equal parts using a cutting plane orthogonal to a coordinate axis.
 - Recursively cut the resulting subdomains.





Nodal Coordinates — RCB Advantages

- Conceptually simple; fast and inexpensive.
- Regular subdomains.
 - Can be used for structured or unstructured applications.
- Effective when connectivity info is not available.
- **Incremental, but no control** of communication costs.








Nodal Coordinates — Coordinate Bisection

- Partition the domain along hyperplanes with node coordinates



•



Good choice of coordinate system leads to inertial bisection

Nodal Coordinates — Inertial Partitioning

- Choose a line L , and then choose a line H orthogonal to it, with half the nodes on either side

1) Center of mass: x_m, y_m

(2) Choose a line L through the points:

L given by $a*(x-x_m)+b*(y-y_m)=0$
with $a^2+b^2=1$; (a, b) is a unit
vector orthogonal to L

(3) Project each point to the line


For each $n_j = (x_j, y_j)$ compute coordinate
 $S_j = -b*(x_j-x_m) + a*(y_j-y_m)$ along L

(4) Compute the median

– Let $S_k = \text{median}(S_1, \dots, S_n)$


(5) Use median to partition the nodes

– Let nodes with $S_j < S_m$ be in V_1 , rest in V_2



Nodal Coordinates — Inertial Partitioning, Choosing L


- Clearly prefer L on left below



- Mathematically, choose L to be a **total least squares fit of the nodes**
 - Minimize sum of squares of distances to L (green lines on last slide)
 - Equivalent to choosing L as axis of rotation that minimizes the moment of inertia of nodes (unit weights) - source of name

Nodal Coordinates — Inertial Partitioning, Choosing L

- $\sum_j (\text{length of } j\text{-th green line})^2$
 $= \sum_j [(x_j - x_m)^2 + (y_j - y_m)^2 - (-b*(x_j - x_m) + a*(y_j - y_m))^2]$
... Pythagorean Theorem



$$\begin{aligned}
 &= (1 - b^2) * \sum_j (x_j - x_m)^2 + 2*a*b* \sum_j (x_j - x_m)*(y_j - y_m) + (1-a^2) \sum_j (y_j - y_m)^2 \\
 &= a^2 * \sum_j (x_j - x_m)^2 + 2*a*b* \sum_j (x_j - x_m)*(y_j - y_m) + b^2 \sum_j (y_j - y_m)^2 \\
 &= a^2 * \begin{matrix} X_1 \\ X_2 \\ X_3 \end{matrix} + 2*a*b* \begin{matrix} X_1 \\ X_2 \\ X_3 \end{matrix} + b^2 * \begin{matrix} X_1 \\ X_2 \\ X_3 \end{matrix} \\
 &= \begin{bmatrix} a & b \end{bmatrix} * \begin{bmatrix} X_1 & X_2 \\ X_2 & X_3 \end{bmatrix} * \begin{bmatrix} a \\ b \end{bmatrix} = \underline{\text{minimum}} = \lambda
 \end{aligned}$$

Minimizing λ ?

Minimized by choosing

$(x_m, y_m) = (\sum_j x_j, \sum_j y_j) / n = \text{center of mass}$
 (a, b)


$$M u = \lambda u \leftrightarrow$$

$$u^T M u = u^T \lambda u = \lambda u^T u = \lambda$$

Nodal Coordinates — Summary

- Algorithms using nodal coordinates are efficient
- Rely on graphs having nodes connected (mostly) to “nearest neighbors” in space
 - algorithm does **not depend on where actual edges** are!
- Common when graph arises from physical model
- **Ignores edges**, but can be used as good starting guess for subsequent partitioners that do examine edges
- Can do very poorly if graph connection is not spatial

Example (graph that is not spatial connected)



In the printed version, the solutions can be found in the appendix

Content


- Motivation for graph partitioning
 - Overview of heuristics
 - Partitioning with nodal coordinates
 - Ex: In finite element models, node at point in (x,y,z) space
 - **Partitioning without Nodal Coordinates**
 - **Ex: In model of WWW, nodes are web pages**
- Spectral Methods**

Coordinate-Free — Spectral Methods

- **Spectral methods** as an example for **global partitioning** algorithms
- Heavily use of Eigenvalue/Eigenvector analysis


Coordinate-Free — Spectral Methods

- Based on theory of Fiedler (1970s), popularized by Horst Simon (1995)
- First motivation with vibrating string
- Label nodes by whether mode - or + to partition into V_1 and V_2



Coordinate-Free — Spectral Methods, Basic Definitions

- **Definition:** The **Laplacian matrix $L(G)$** of a graph $G(V, E)$ is a $|V|$ by $|V|$ symmetric matrix, with one row and column for each node. It is defined by
 - $L(G) (i,i) = \text{degree of node } i$ (number of incident edges)
 - $L(G) (i,j) = -1$ if $i \neq j$ and there is an edge (i,j)
 - $L(G) (i,j) = 0$ otherwise



$$\begin{array}{c}
 \begin{matrix} 1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9 \end{matrix} \\
 \begin{matrix} 1 \\ 2 \\ 3 \\ 4 \\ 5 \\ 6 \\ 7 \\ 8 \\ 9 \end{matrix} \left[\begin{array}{ccccccccc}
 2 & -1 & & & & & & & \\
 -1 & 3 & -1 & & & & & & \\
 & -1 & 2 & & & -1 & & & \\
 -1 & & & 3 & -1 & & -1 & & \\
 & & -1 & -1 & 4 & -1 & & -1 & \\
 & & & -1 & -1 & 3 & & & -1 \\
 & & & & -1 & & 2 & -1 & \\
 & & & & & -1 & -1 & 3 & -1 \\
 & & & & & & -1 & -1 & 2
 \end{array} \right]
 \end{array}$$

Properties of Laplacian matrices

- **Theorem**: Given a graph G , $L(G)$ has the following properties
 - $L(G)$ is symmetric — this means the eigenvalues of $L(G)$ are **real** and its **eigenvectors are real** and **orthogonal**.
 - Let $e = [1, \dots, 1]^T$, i.e. the column vector of all ones. Then $L(G)e = 0$
 - The eigenvalues of $L(G)$ are **nonnegative**:
 $0 = \lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_n$
 - The number of connected components of G is equal to the number of λ_i equal to 0.
- **Definition**: $\lambda_2(L(G))$ is the **algebraic connectivity** of G
 - The magnitude of λ_2 measures connectivity
 - In particular, $\lambda_2 \neq 0$ if and only if G is connected

Relation between Laplace Matrix and Graph Partitioning

- **Theorem (Fiedler, 1975):**

Let G be connected, $L(G)$ the Laplace matrix, and N_+ and N_- a partitioning with

$$x(i) = +1 \quad \text{if } v_i \text{ in } N_+$$

$$x(i) = -1 \quad \text{if } v_i \text{ in } N_-.$$

Then we have the following property:

$$\begin{aligned} \text{\#edge-cut between } N_+ \text{ and } N_- \\ = \quad \frac{1}{4} * \mathbf{x}^T * \mathbf{L}(G) * \mathbf{x} \end{aligned}$$

Proof: (next slide)

Relation between Laplace Matrix and Graph Partitioning

$$\begin{aligned}
 x^T \cdot L(G) \cdot x &= \sum_j \sum_i L(G)_{(i,j)} \cdot x_i \cdot x_j \\
 &= \sum_{i=j} L(G)_{(i,i)} \cdot x_i^2 + \sum_{i \neq j} L(G)_{(i,j)} \cdot x_i \cdot x_j \\
 &= \sum_{i=j} L(G)_{(i,i)} \cdot x_i^2 \\
 &\quad + \sum_{i \neq j; i, j \in N^+} L(G)_{(i,j)} \cdot x_i \cdot x_j + \sum_{i \neq j; i, j \in N^-} L(G)_{(i,j)} \cdot x_i \cdot x_j \\
 &\quad + \sum_{i \neq j; i \in N^+, j \in N^-} L(G)_{(i,j)} L(G)_{(i,j)} \cdot x_i \cdot x_j \\
 &= \sum_i \text{degree}(i) \\
 &\quad + \sum_{i \neq j; i, j \in N^+} (-1) \cdot (+1) \cdot (+1) + \sum_{i \neq j; i, j \in N^-} (-1) \cdot (-1) \cdot (-1) \\
 &\quad + \sum_{i \neq j; i \in N^+, j \in N^-} (-1) \cdot (+1) \cdot (-1)
 \end{aligned}$$

Relation between Laplace Matrix and Graph Partitioning

$$\begin{aligned} x^T \cdot L(G) \cdot x &= \sum_{i,j} L(G)_{(i,j)} \cdot x_i \cdot x_j \\ &= \sum_i \text{degree}(i) \\ &\quad + \sum_{i \neq j; i, j \in N^+} (-1) \cdot (+1) \cdot (+1) + \sum_{i \neq j; i, j \in N^-} (-1) \cdot (-1) \cdot (-1) \\ &\quad + \sum_{i \neq j; i \in N^+, j \in N^-} (-1) \cdot (+1) \cdot (-1) \\ &= 2 \cdot \# \text{edges in } G \\ &\quad - 2 \cdot (\# \text{edges connecting node in } N^+ \text{ to nodes in } N^+) \\ &\quad - 2 \cdot (\# \text{edges connecting node in } N^- \text{ to nodes in } N^-) \\ &\quad + 2 \cdot (\# \text{edges connecting node in } N^+ \text{ to nodes in } N^-) \\ &= 4 \cdot (\# \text{edges connecting node in } N^+ \text{ to nodes in } N^-) \end{aligned}$$

Relation between Laplace Matrix and Graph Partitioning

- With the theorem we can formulate the **graph bisection** as a discrete optimization problem

$$\begin{array}{ll} 1. & |V_1| = |V_2| \quad \Leftrightarrow \sum_i x(i) = 0 \\ 2. & \min \text{ #cut edges between } V_1 \text{ and } V_2 \quad \Leftrightarrow \min x^T * L(G) * x \end{array}$$

or

$$\begin{array}{ll} \min & f(x) = \frac{1}{4} x^T * L(G) * x \\ \text{constraints} & x_i = \{+/- 1\}, \quad x^T * x = n \\ & x^T * e = 0 \text{ with } e = [1, 1, \dots, 1]^T \end{array}$$

- The **discrete combinatorial** problem is NP-hard \rightarrow use a **continuous problem**

$$\begin{array}{ll} \min & f(z) = \frac{1}{4} z^T * L(G) * z \\ \text{constraints} & z^T * z = n, \text{ } z \text{ real vector} \\ & z^T * e = 0 \text{ with } e = [1, 1, \dots, 1]^T \end{array}$$

Relation between Laplace Matrix and Graph Partitioning

- Let's try to solve the continuous graph bisection problem

Relation between Laplace Matrix and Graph Partitioning

- Minimal solution of $z^T * L(G) * z$ is easy to find.
- $L(G)$ is symmetric $\rightarrow L(G)$ has n orthonormal eigenvectors u_1, \dots, u_n with eigenvalues $0 = \lambda_1 \leq \dots \leq \lambda_n$ and $u_1 = \frac{1}{\sqrt{n}} * e$, $e = [1, 1, \dots, 1]^T$.
- A vector z is a linear combination of eigenvectors u_i :

$$z = \sum \alpha_i u_i = \alpha_1 u_1 + \alpha_2 u_2 + \dots + \alpha_n u_n.$$
- **First constrained:** $z^T * e = 0$ or $z^T * u_1 = 0$ it is necessary that

$$z^T * u_1 = (\sum \alpha_i u_i)^T * u_1 = \alpha_1 u_1^T * u_1 = \alpha_1 \Rightarrow \alpha_1 = 0$$
- **Second constrained:** $z^T * z = n$ it is necessary that

$$z^T * z = (\sum \alpha_i u_i)^T * (\sum \alpha_j u_j) = \sum \sum \alpha_i \alpha_j u_i^T * u_j = \sum \alpha_i^2 = n$$
- **Minimize $4 * f(z) = z^T * L(G) * z$**

$$z^T * L(G) * z = (\sum \alpha_i u_i)^T * L * (\sum \alpha_j u_j) = (\sum \alpha_i u_i)^T * (\sum \alpha_j \lambda_j u_j) = \sum \sum \alpha_j \alpha_i \lambda_j u_i^T * u_j = \sum \alpha_i^2 \lambda_j \geq \lambda_2 \sum \alpha_i^2 = \lambda_2 * n$$

Relation between Laplace Matrix and Graph Partitioning

- **Minimize** $4 * f(z) = z^T * L(G) * z$

$$\begin{aligned} z^T * L(G) * z &= (\sum \alpha_i u_i) * L * (\sum \alpha_j u_j) = (\sum \alpha_i u_i)^T * (\sum \alpha_j \lambda_j u_j) = \\ \sum \sum \alpha_j \alpha_i \lambda_j u_i^T * u_j &= \sum \alpha_i^2 \lambda_j \geq \lambda_2 \sum \alpha_i^2 = \lambda_2 * n \end{aligned}$$

- Minimum is at $z = \sqrt{n} * u_2$.

- **Spectral Bisection Algorithm:**

- Compute eigenvector u_2 corresponding to $\lambda_2 (L(G))$
- For each vertex v of G
 - if $u_2(v) < 0$ put node v in partition V_1
 - else put vertex v in partition V_2

- The second eigenvector u_2 is called **Fiedler Eigenvector** of the Graph Partitioning problem.

Demo – Partitioning in Matlab and Julia

